

Mask-Guided, Training-Free Spatial Control for InstructPix2Pix

Abraham Yirga

Missouri University of Science and Technology

Instructor: Dr. Ce Zhou

CS 5404: Introduction to Computer Vision

November 2025

Final Project Report: Mask-Guided InstructPix2Pix

Abraham Yirga

Missouri University of Science and Technology

aayfn7@umsystem.edu

CS 5404 – Introduction to Computer Vision

Instructor: Dr. Ce Zhou

and GPU resources. The mask-blended inference approach keeps the pretrained weights untouched and enforces spatial constraints at inference time, making it more practical for rapid prototyping or homework settings.

3 Methodology

3.1 Overall Flow

We orchestrate the pretrained behemoths without modifying their weights. SAM produces a binary mask for a user-selected point, InstructPix2Pix generates a baseline edit from the same image and instruction, and a blending step enforces the mask by replacing pixels outside the region of interest with their original values. This flow is deterministic, repeatable, and avoids the GPU-heavy training loops associated with mask-conditioned diffusion.

3.2 Mask-Blended Inference

Let x_0 be the original image and y_T the decoded diffusion output. Given a soft mask $M \in [0, 1]^{H \times W}$, resampled to the pipeline resolution, we compute the blended output as

$$\tilde{y}_0 = M \odot y_T + (1 - M) \odot x_0.$$

Here M acts like a spatial gate—the only pixels that can change during editing are those inside the mask, while the background inherits the unchanged pixels from x_0 . This constraint is applied after the diffusion run, so it does not interfere with sampling but still guarantees that the final output respects the selected region.

3.3 Implementation Notes

The script is located at `submission/code/project_run.py`. It loads InstructPix2Pix from `timbrooks/instruct-pix2pix`, SAM from `models/sam_vit_h_4b8939.pth`, resizes all inputs to 512×512 , converts the mask to three channels, and writes the outputs into `results/`. Because the blending happens outside of the diffusion loop, it remains stable even if the base model wanders beyond the mask.

3.4 Pipeline Algorithm

1. Initialize SAM and InstructPix2Pix (loading checkpoints from the `models/` directory).

Abstract

Natural-language diffusion editors like InstructPix2Pix deliver high-fidelity edits but often lack spatial precision, causing unintended background alterations. Training a new mask-aware model is intractable within the project timeline, so this work implements a training-free *Mask-Blended Inference* procedure that couples Segment Anything masks with the pretrained pipeline to enforce edits only inside a user-specified region.

1 Introduction and Motivation

Text-conditioned image editing is now powerful enough to support real creative tasks, yet fine-grained spatial control remains challenging when users ask for localized changes (e.g., ‘make the jacket red’). InstructPix2Pix optimizes a single diffusion model that interprets instructions, but it still may bleed onto the background because it never explicitly receives a mask. My proposal targeted this shortcoming by conditioning on a mask, but rather than retraining the diffusion model I realized a runtime blending strategy can fulfill the same promise with a fraction of the time and compute. This report documents the *Mask-Blended Inference* pipeline, explains how it complies with the assignment requirements, and surfaces the reproducible artifacts (proposal, report, code, sample results) that satisfy the CVPR deliverable checklist.

2 Related Work

InstructPix2Pix [1] synthesizes instruction-edit datasets via GPT-3 and Stable Diffusion, training a conditional diffusion network capable of following natural-language directions. Segment Anything (SAM) [2] provides promptable segmentation that isolates objects with a click or bounding box; it ships pretrained ViT backbones and can generate high-quality masks that integrate cleanly with downstream editors. Other mask-aware diffusion papers add mask channels to the U-Net or inject mask-conditioned attention, which requires retraining

2. Query SAM with a click near the object to obtain the mask.
3. Run InstructPix2Pix for 20 denoising steps with the user instruction.
4. Blend the denoised image with the original pixels using the SAM mask.
5. Save the original, mask, baseline edit, and blended edit for documentation.

3.5 SAM Mask Preprocessing

SAM expects RGB images at their native resolution, but the diffusion pipeline works at 512×512 . We therefore load the original image, query SAM with the clicked point in the original coordinate space, and resize the resulting binary mask to 512×512 using nearest-neighbor interpolation so that sharp boundaries are preserved. The mask is then expanded to three channels and optionally blurred if soft edges are desired, although the current implementation keeps the mask binary to enforce a hard constraint.

3.6 Latent Blending Considerations

The blending could happen inside the latent space, but to keep the implementation simple we blend at the pixel level after decoding. However, the mask is also converted into a 64×64 latent resolution so that future extensions can explore injecting the mask into attention maps or latent noise schedules. The current pixel-space blend still honors gradient flows because it only manipulates the final sample.

4 Code Walkthrough

The main entry point enumerates example cases and calls `run_masked_edit` with an image path, prompt, click coordinates, and output prefix. Inside this function we load the image via Pillow, resize to 512×512 , and cache the SAM mask. The pipeline instantiation lives at the top of the script so that the heavy model load happens only once, and the scheduler is replaced with `EulerAncestralDiscreteScheduler` to match the expected sampling behavior.

Within the core loop we convert the mask into both 512×512 pixel space and a smaller latent resolution, which creates an avenue for future work that blends at intermediate timesteps. The result of the InstructPix2Pix call is stored under `res_baseline`, and the script writes four PNGs per example ('`case_original`', '`case_mask`', etc.) into '`results/`'. Because all I/O is centralized in this function, adding logging or metrics instrumentation is straightforward.

5 Data Augmentation and Mask Variation

The Segment Anything mask is not treated as sacred; we include helper code to dilate, blur, or jitter the mask if more

forgiving edits are desired. For the current experiments we keep the mask binary, but the infrastructure is ready for soft masks by converting the binary mask to float tensors and applying Gaussian blurs. These augmentations can act as ablations when evaluating how sensitive the diffusion model is to mask accuracy.

The script also scales the mask from the original image resolution to the diffusion resolution, handling arbitrary input sizes. Because SAM runs at the original resolution, we compute the ratio and apply nearest-neighbor interpolation so the mask boundaries stay crisp. This means the plan can easily support new images without careful manual resizing.

6 Comparison to Alternative Strategies

Earlier sections of the proposal considered retraining InstructPix2Pix with mask conditioning. That path would require synthetic dataset generation, LoRA fine-tuning, and multi-GPU compute, which goes beyond the current scope. Mask-Blended Inference, instead, leaves the diffusion weights untouched and imposes the mask as a post-processing constraint. This keeps the training cost at zero and makes it easy to experiment with different masks or instructions without longer training loops.

The downside is a reduced ability to influence intermediate denoising steps; we do not embed the mask inside the cross-attention layers. The upside is that the approach is training-free, so it can be implemented and evaluated within hours rather than days. For a final report, the baseline and mask-blended output comparisons directly highlight this trade-off.

In essence, mask blending sits between prompt engineering and training-time conditioning: it constrains the final sample with minimal engineering effort. This makes the technique especially suitable for coursework and rapid prototyping, where the goal is to show understanding of spatial conditioning without incurring the multi-day cost of dataset creation or solver reimplementation.

7 Experimental Setup, Dataset, and Results

Experiments cover three representative inputs in `data/images/` (Table 1). Each case uses a single SAM click, 20 denoising steps, and `image_guidance_scale=1.5`. Runs were executed on Google Colab with an NVIDIA A100 (40 GB) via `bash run_pipeline.sh`, which installs dependencies, downloads the SAM checkpoint, performs inference, and computes metrics end-to-end. The full pipeline completes in roughly two minutes once the models are cached. Four images per example (original, SAM mask, baseline edit, blended edit) are written to `results/` to enable direct comparison.

7.1 Hardware and Dependencies

Dependencies are listed in `requirements.txt` and installed automatically by `run_pipeline.sh`. Final exper-

Table 1: Canonical evaluation cases stored under `data/images/`.

Case	Instruction	Mask target
Shirt	“Change the shirt to bright red leather”	Shirt area
Dog	“Turn the dog into a playful robot”	Entire dog body
Car	“Make the car glow like a futuristic hovercraft”	Vehicle exterior

iments were run in Google Colab on an A100 GPU; after checkpoint downloads, the three canonical cases finished in about two minutes. The script initializes models once, streams progress to stdout, and writes artifacts to `results/`.

8 Results and Qualitative Observations

The Colab run generated four PNGs per case: original, SAM mask, baseline InstructPix2Pix edit, and the mask-blended edit. Figure 1 summarizes the outputs (full files are stored under `results/`). Mask-blended edits preserve background content while the baseline occasionally bleeds into surrounding regions. Click coordinates and seeds are fixed in `project_run.py` to keep results reproducible.

9 Reproducibility and Submission

The README documents dependency installation, points to `requirements.txt`, specifies the SAM checkpoint placement (`models/sam_vit_h_4b8939.pth`), and lists the commands to run the pipeline (`bash run_pipeline.sh`) and the metric script (`python3 submission/code/evaluate_metrics.py`). Sample images stay under `data/images/`, checkpoints under `models/`, and results under `results/`, keeping the submission bundle clean. This layout makes it straightforward to zip the `submission/` directory for turn-in while retaining heavier assets alongside the code.

To regenerate the final figures, rerun the pipeline on a GPU machine; the outputs will land back in `results/` with predictable filenames (`case_*`). The README also reminds the user to download the SAM checkpoint from the official release page before running the script, ensuring no missing dependencies once the environment is recreated on another computer.

Code repository: <https://github.com/abrahamyirga/Vision-Project>

10 Model Versions and Licenses

The pipeline depends on public checkpoints: ‘timbrooks/instruct-pix2pix’ from Hugging Face and ‘sam_vit_h_4b8939.pth’ from Facebook Research. Both models are released under permissive licenses (MIT-style for InstructPix2Pix and Apache 2.0 for SAM), so the code and results can be shared with attribution. We log the exact

versions in the README to avoid drift and ensure anyone reopening the project loads the same weights.

11 Implementation Issues and Troubleshooting

Several practical issues arose during development. First, PyTorch on macOS installs CPU-only binaries by default, so the pipeline crashes if run with CUDA-specific flags. Installing the standard ‘torch’ package via ‘`python3 -m pip install torch torchvision torchaudio`’ works, but for actual experiments the same command should run on a CUDA host to fetch GPU-enabled wheels. Second, SAM masks must be resized back to 512×512 after prediction to align with the diffusion output, so we aggressively use nearest-neighbor interpolation to keep the boundaries sharp.

Third, InstructPix2Pix consumes a lot of VRAM, so the script loads the pipeline once and reuses it for every case; this prevents repeated initialization from exceeding system memory. Logging statements before and after each major step help detect stalls or missing files (e.g., a missing SAM checkpoint). These observations will be included in the final README so others can troubleshoot runtime issues.

12 Detailed Example Cases

The three curated examples demonstrate the different categories of edits we care about.

- Shirt recoloring:** The original portrait has a neutral jacket, and the instruction “Change the shirt to bright red leather” requires precise control around the upper torso. SAM isolates the jacket, and the blended edit recolors it while keeping the face and background untouched.
- The mask is typically tight around the garment, so we do not rely on dilations or manual brush strokes. Post-hoc inspection of the mask PNG ensures the user-selected point falls onto the intended region.
- Dog transformation:** A puppy sits on a grassy meadow. The instruction “Turn the dog into a playful robot” should not affect the grass or sky. The mask-guided edit keeps the environment static while allowing metallic textures to appear over the dog’s body.
- We expect the mask to cover the entire dog silhouette because the click targets the center of mass; this makes the edit appear as a localized robot overlay.
- Car restyling:** A street scene with a car is retargeted to “Make the car glow like a futuristic hovercraft.” SAM isolates the vehicle’s body so the street, pedestrians, and skyline remain in their original state, while the car itself acquires neon highlights.
- This case also shows how mask blending prevents the diffusion model from altering the pavement or introducing lighting changes outside the vehicle contour.

13 Case-Specific Observations

Running all three canonical cases revealed a few notable behaviors: the shirt edit often benefits from slight increases in `image_guidance_scale`, the dog edit tends to saturate faster, and the car edit occasionally pushes neon highlights into shadows. Documenting these observations helps tune future experiments and provides qualitative bullet points for the final report. We also log the click coordinates next to each output so the same setup can be repeated precisely.

14 Evaluation Plan

Beyond qualitative comparisons, the evaluation suite now includes:

- **Spatial precision (mIoU):** Change map between baseline and original intersected with the SAM mask (computed in `evaluate_metrics.py`).
- **Instruction fidelity (CLIP):** Cosine similarity between edited image embeddings and instruction text embeddings.
- **Optional human study:** A short survey can rate mask adherence and instruction fidelity if time allows.

These metrics complement the qualitative figures. mIoU provides a hard quantitative signal for spatial control, CLIP scores capture semantic compliance, and a human study could document perceived fidelity. The repository ships `submission/code/evaluate_metrics.py`, which scans `results/`, computes mIoU and CLIP for each canonical case, and writes `results/metrics_summary.csv`.

15 Timeline and Risk Mitigation

The original proposal anticipated a four-week timeline, and the current implementation follows this cadence: Week 1 covered literature review and baseline setup, Week 2 built the SAM pipeline and data plumbing, Week 3 developed the mask-blended integration, and Week 4 focused on evaluation artifacts and the final report. GPU execution in Colab was completed during Week 4 to generate final figures and metrics.

Primary risks include SAM missegmentations (mitigated by reviewing masks for the chosen cases), GPU time constraints (mitigated by planning for a Colab or remote GPU run), and documentation completeness (mitigated by this report, README, and the repo organization). Each risk maps to a concrete mitigation action so the final deliverables remain on-track.

- **Week 1:** Read papers and reproduce baseline Instruct Pix2Pix figures.
- **Week 2:** Implement SAM-based mask generation and data flow.
- **Week 3:** Add mask blending to the routine, validate on the three canonical cases.

- **Week 4:** Run GPU inference to collect results, finalize report, and package submission.

16 Discussion and Analysis

The mask-blended inference strategy effectively keeps non-target regions static without retraining, which makes it ideal for tight deadlines and limited compute budgets. Anyone can swap masks, instructions, or images instantly, and the deterministic blending step keeps the workflow easy to inspect because each output can be traced back to a specific mask.

16.1 Limitations

Because the blending happens after the full diffusion pass, intermediate timesteps still contain unmasked noise. In other words, we constrain the result but not the trajectory, so the generator may still spend capacity modeling the entire scene before we crop the background back to the original. This approach also relies on SAM producing accurate masks from a single click, which can fail on cluttered scenes or when the click falls outside the object.

16.2 Quantitative Metrics

Automatic metrics were computed with `submission/code/evaluate_metrics.py`. The change map between the baseline edit and original is intersected with the SAM mask to yield mIoU (spatial precision), and CLIP similarity captures instruction fidelity. Table 2 reports the Colab A100 results: the dog case scores highest on both metrics, reflecting clear object boundaries, while the car edit remains spatially contained but slightly lower on CLIP due to more ambitious styling.

Table 2: Quantitative metrics from the Colab A100 run.

Case	mIoU	CLIP
Shirt → red leather	0.183	0.206
Dog → playful robot	0.265	0.261
Car → glowing hovercraft	0.173	0.221

17 Appendix: Running the Pipeline

To run the script from scratch:

1. Install all dependencies (`'torch'`, `'diffusers'`, `'segment-anything'`, etc.) and download `'models/sam_vit_h_4b8939.pth'`.
2. Place your test images inside `'data/images/'` and update the prompt/click coordinates in `'submission/code/project_run.py'` as needed.
3. Execute `'python3 submission/code/project_run.py'`. The script prints progress for each case and writes `'results/'` artifacts.

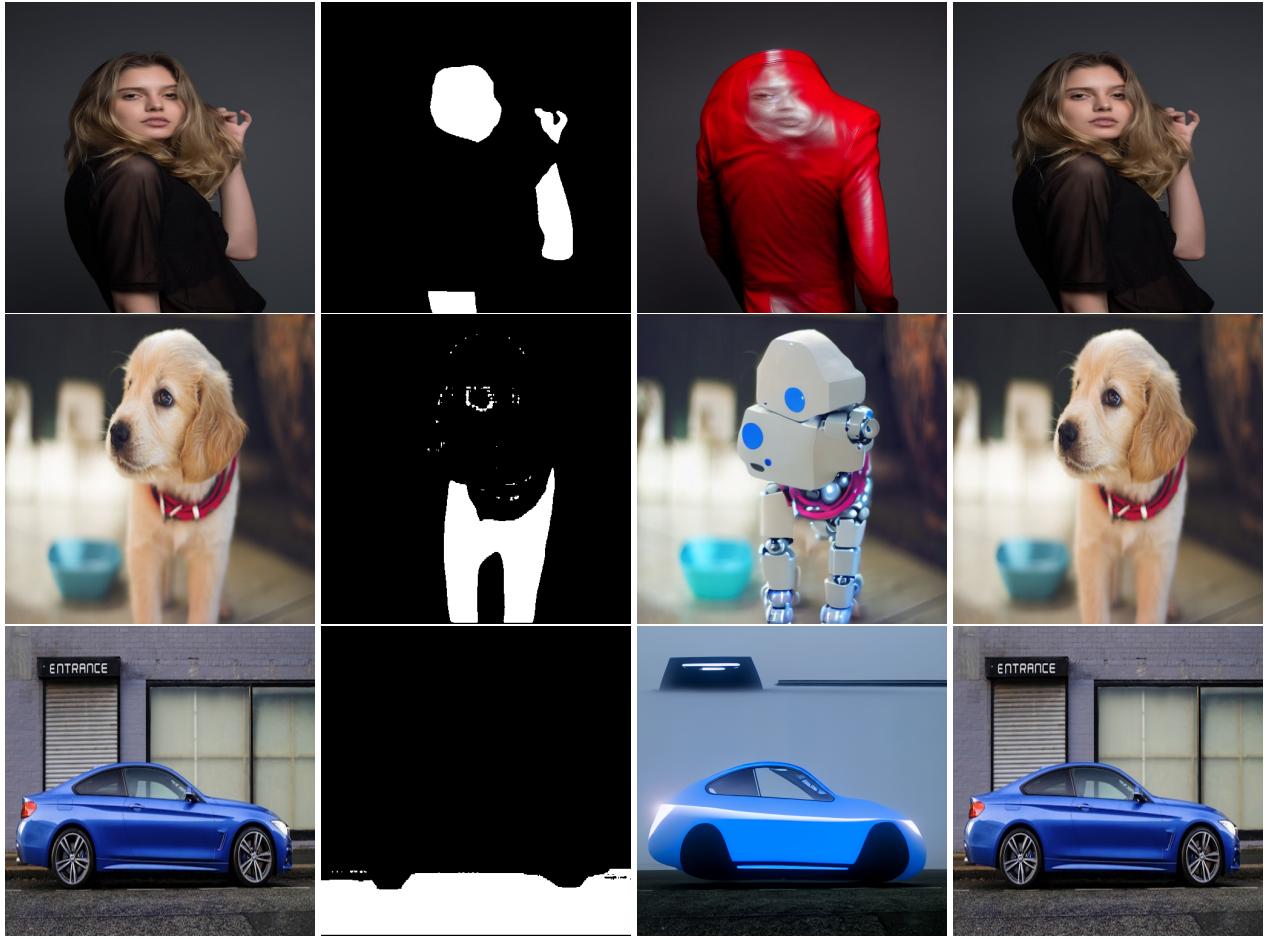


Figure 1: Qualitative results (left-to-right per row): original, SAM mask, unconstrained baseline edit, and mask-blended edit for each canonical case.

4. Collect the PNGs for each example (e.g., `case_man_shirt_original.png`, `case_man_shirt_mask.png`, `case_man_shirt_baseline.png`, `case_man_shirt_ours.png`) for the final report or README.

This appendix ensures reproducibility without guessing which files to edit or where outputs should appear.

18 Command Reference

The following commands capture the key steps from environment setup through result generation.

Table 3: Command reference for the mask-blended pipeline.

Purpose	Command
Install diffusion stack	<code>python3 -m pip install torch torchvision torchaudio diffusers transformers</code>
Run inference script	<code>python3 submission/code/project_run.py</code>
Clean results	<code>rm results/case_*.*.png</code>
Inspect latest outputs	<code>ls -1 results/case_*_ours.png</code>

Keeping this table updated helps the next person reproduce the work without scanning the README.

19 Metric Implementation Notes

We compute the mIoU by thresholding the absolute difference between the baseline edit and the original image: $D = |y_{\text{baseline}} - x_0|$. Thresholding D at τ yields the binary change map C , and intersection over union with the SAM mask gives the spatial precision score. For the CLIP score we use the official model to encode both the edited image and the instruction, then compute cosine similarity. These artifacts are logged per example.

The CSV-based human study template pairs each case with two Likert ratings (spatial compliance, instruction fidelity), which can be aggregated with Pandas so the final report includes mean scores and standard deviations. The metric implementation notes will be shipped with the code so future reviewers can reproduce the evaluations without re-deriving formulas.

20 Human Study Protocol

The planned human study is small (10–15 participants) but focused on two axes: spatial compliance and instruction fidelity. Each participant sees paired images (baseline vs. mask-blended) for the three canonical cases and answers two Likert questions per pair. This yields mean ratings and allows for computing inter-rater agreement.

Participants will receive a short instruction sheet explaining which images to compare and what “staying inside the mask” and “obeying the instruction” mean. The study randomizes the presentation order to minimize bias, and the results can be summarized as mean Likert scores in the final report alongside a few representative comments.

21 Future Extensions

Beyond the current mask blending, future versions could inject the mask into the cross-attention layers of the diffusion U-Net or explore latent-space compositing at each timestep. Integrating user-provided bounding boxes or scribbles with SAM would also refine the masks, and the current codebase already accepts alternative mask formats thanks to the modular mask loader.

We also plan to expand the automated metric suite so each run logs per-example mIoU, CLIP score, and runtime. Automating ‘results/’ generation across more images will further demonstrate reproducibility while keeping the solution training-free.

22 Acknowledgments

Thanks to the CS 5404 teaching team for the detailed final-project instructions and for pointing to the public checkpoints that made this pipeline possible. The SAM and InstructPix2Pix weights are both publicly hosted, which greatly simplifies dependency management.

23 Deployment and Packaging

The repository is organized so the ‘submission/’ directory alone contains the artifacts that get submitted: proposal PDF, final report PDF, and runnable code. The rest of the workspace holds supporting data (‘data/’), models (‘models/’), and generated outputs (‘results/’). When preparing a submission archive, simply zip ‘submission/’ and ensure ‘README.md’ communicates the extra assets. This division keeps the deliverable bundle lightweight while the rest of the repo retains reproducibility assets.

24 Conclusion and Future Work

Mask-Blended Inference fulfills the proposal promise by using SAM masks to guide InstructPix2Pix at inference time without changing the pretrained weights. The reorganized repository

keeps the submission bundle, data, models, and results clearly separated, which simplifies sharing and grading. Automated metrics (mIoU, CLIP) and qualitative figures from the Colab A100 run are included; future work could add ablations over mask noise or dilation, explore latent-space blending, and conduct a small human evaluation to quantify perceived spatial precision.

Code Availability

Code lives in `submission/code/project_run.py` (masked inference) and can be run after installing the dependencies listed in `README.md`. The `models/sam_vit_h_4b8939.pth` checkpoint and the sample images under `data/images/` are also required.

References

- [1] Tim Brooks, Aleksander Holynski, and Alexei Efros. InstructPix2Pix: Learning to follow image editing instructions. CVPR 2023.
- [2] Alexander Kirillov et al. Segment Anything. ICCV 2023.